

Paras Gupta

+91 8127683667 parasg484@gmail.com [LinkedIn](#)

WORK EXPERIENCE

Data Engineer 2 (Client- HDFC Bank)

NexGen Tech, Noida, INDIA

June 2025 - Present

- Led end-to-end migration of **HDFC loan analytics** from **Azure Data Lake** → **Lakehouse**, building **Apache Iceberg tables (Trino)** and orchestrating pipelines via Python-based **Airflow DAGs**.
- Built scalable **batch and near real-time ingestion pipelines** from **REST APIs**, **Adobe Analytics API**, **Kibana/Elasticsearch API**, **on-prem PostgreSQL**, and **SFTP** using **incremental loading**, **schema evolution handling**, and **metadata-driven processing**.
- Contributing to real-time pipelines using **Kafka**, **Pub/Sub** and **Azure Data Factory**, processing **20M+** daily loan clickstreams with low latency.
- Leveraged **Unity Catalog** to enforce fine-grained **data governance**, **access controls**, and **lineage tracking**, enabling seamless collaboration across Data Architecture, BI, and QA teams.
- Designed **Python-based UDFs** in **Databricks** to parse **complex JSON, XML, and HTML payloads**, reducing manual intervention by **90%**.
- Optimized Databricks clusters, Delta tables (**partitioning**, **Z-Ordering**), and ADF pipelines — achieving **20%** cost savings and **30%** faster runtimes.
- Architected a **metadata-driven ingestion & transformation framework** using **PySpark**, **Delta Lake**, **ADF**, and **Databricks Workflows**, increasing pipeline reusability by **50%**.

Data Engineer 1 (Client- Microsoft)

MAQ Software, Noida, INDIA

January 2023 - May 2025

- Enhanced a **Medallion Architecture** based **3-layer Delta Lakehouse (Bronze → Silver → Gold / Staging → Curated → Service)** in **Azure Databricks** to improve scalability, **data quality**, **governance**, **security**, and **analytics readiness**.
- Reduced **Azure Databricks** notebook performance overhead by **40%** via advanced **caching**, efficient Spark configurations, and **partitioning techniques**.
- Developed high-performance **PySpark and SQL data transformation pipelines** with Spark optimization techniques including **caching**, **partition pruning**, and **optimized joins**, improving client decision-making by **15%**.
- Developed **Power BI** dashboards, reducing report refresh times by **25%** through **DAX optimization** and efficient data modeling.

PROJECTS

Personal/Business Loan Engine Analytics | NexGen Tech / HDFC Bank

- Led migration from **Azure Data Lake** → **Lakehouse**; built **Apache Iceberg tables (Trino)** and orchestrated pipelines via **Airflow DAGs** ingesting from **REST APIs**, **Adobe Analytics API**, **Elasticsearch**, and **SFTP**.
- Designed **Star Schema** data models and transformed complex **JSON/XML** loan data using **PySpark** and **SQL**; built custom Python parsers for analytics-ready datasets processing **20M+** daily events ensuring scalable **ETL/ELT processing**.
- Applied **Unity Catalog** for data governance and lineage; implemented near real-time streaming using **Kafka** and **Pub/Sub** with **20%** cost savings through cluster and Delta table optimizations.

Microsoft Modern Work Sales Engine Analytics | MAQ Software / Microsoft

- Architected and enhanced a **Medallion Architecture based 3-layer Delta Lakehouse (Bronze → Silver → Gold)** in **Azure Databricks**; reduced notebook overhead by **40%** through advanced caching, partitioning, and Spark tuning.
- Developed high-performance **PySpark and SQL** pipelines to generate sales KPIs and funnel metrics, directly influencing a **15%** improvement in sales decision-making for the Modern Work business unit.
- Built **Power BI** dashboards with **DAX optimization** and efficient data modeling, reducing report refresh times by **25%** and enabling self-serve analytics for sales and operations teams.

EDUCATION

PSIT College of Engineering, Kanpur

B.Tech, Computer Science & Engineering

Jul 2019 - Jun 2023

8.86 CGPA

ST. Thomas College, Lucknow

Senior Secondary Certificate

Apr 2018 - Mar 2019

92%

TECHNICAL SKILLS

Languages: Python, SQL, PySpark, DAX, KQL, Bash

Lakehouse: Apache Iceberg (Trino), Delta Lake, ADLS Gen2, Azure Blob Storage, Unity Catalog, Data Lakehouse, Medallion Architecture

Orchestration: Apache Airflow (Python DAGs), Azure Data Factory, Databricks Workflows, Workflow Orchestration, ETL/ELT Pipelines

Streaming: Apache Kafka, Pub/Sub

Cloud - Azure: Azure Databricks, Azure Data Factory, Azure Key Vault, Logic Apps

Data Sources & APIs: REST APIs, Adobe Analytics API, Elasticsearch/Kibana API, SFTP, PostgreSQL, SQL Server

Reporting: Power BI, DAX Optimization

DevOps: Git, Azure DevOps, CI/CD Pipelines, Deployment Automation, Version Control

ACHIEVEMENTS AND CERTIFICATIONS

- LeetCode Profile — Solved 250+ questions.
- HackerRank Profile — Solved 100+ questions.
- Achieved an 83.12 percentile in TCS NQT Cognitive Skills.
- Microsoft Fabric Certified Data Engineer.